

Time Range Queries for Hereditary Properties

6.854 Final Project

Arsen Mamikonyan, Hayk Saribekyan

December 14, 2016

Abstract

Time range queries are important for analysing large datasets with timestamps. Such datasets are common in computational geometry, where a sequence of points with timestamps often corresponds to a trajectory of an object. Time range queries in this case check a certain property of section of the trajectory.

In this paper we review recent general frameworks that can be used to handle property testing queries for time ranges, and discuss their implementations for solving certain problems in computational geometry. The results are limited to *hereditary* properties, which cover a large variety of interesting problems.

For two different hereditary properties, we compared the performance of efficient an algorithm and a naive one in practice. We observed that the efficient algorithms require more than ten times the amount of code in certain cases, but the large constant factor associated with it is compensated when working with large datasets.

1 Introduction

TODO: add normal references (maybe to envelopes), see the first paper, where in the first page there are a lot of motivational references.

With the abundance of GPS and other movement tracking sensors there is a large amount of timestamped location data. The movement of an object in space can be modelled by a sequence of points $S = s_1, \dots, s_n$. To analyse the trajectory of the movement, one may want to check a certain property P for portions $S[i, j] = s_i, \dots, s_j$ of S for given values of i and j . This paper discusses how to efficiently address such queries when P satisfies the following two properties:

- P is boolean.
- P is hereditary. This means that for a given sequence S , if $P(S)$ is true, then $P(S')$ is true for any continuous subsequence S' of S .

In this paper we discuss and implement two properties for a sequence of points $S = s_1, \dots, s_n$:

- *Monotonicity*: S is *monotone* if there is a direction vector v such that all vectors (s_i, s_{i+1}) have a positive projection on v . In terms of trajectories, monotonicity shows whether the object travelled more or less in the same direction.
- *Closeness*: S is *close* if any pair of points in S is at most of distance 1. One could use closeness queries to detect if a moving object mostly stayed in the same surrounding.

For a property P satisfying the restrictions above, let $j^*(i)$ be the largest index j such that P holds for $S[i, j^*(i)] = s_i, \dots, s_{j^*(i)}$. Now notice, that since P is a hereditary property $P(S[i, j]) = \text{true}$ if and only if $i \leq j \leq j^*(i)$. Therefore, if we construct $j^*(i)$ we can answer property testing queries in $O(1)$. Therefore, for the rest of the paper our goal will be to efficiently compute $j^*(i)$.

Bokal et al. [1] propose an algorithms that compute $j^*(i)$ in $O(n)$ time for monotonicity, and in $O(n \log^2 n)$ time for closeness. Chan and Pratt [2] describe a different algorithm to achieve the same result for closeness, but also improve it to $O(n \log n)$ using fractional cascading.

The rest of the paper is organized as follows: Section 2 describes the naive algorithms we developed to solve two problems; in Section 3 we review the general framework that is set in [1] for solving time range query problems, and show how it is applied to compute $j^*(i)$ for monotonicity. Section 4 reviews the framework by [2] and its application for the closeness property. We then give some implementation details in Section 5, and describe the experimental setup and our tests in Section 6.

2 Naive Algorithms

Suppose that for a sequence of length n a property P can be checked in $T_P(n)$ time. Then, one can compute $j^*(i)$ using binary search for each i . On a step of a binary search concerning a range $[l, r]$ the algorithm would spend $T_P(r - l)$ time. The total runtime of this simple algorithm is $O(nT_P(n) \log n)$. Call this algorithm *SN* (stands for super-naive). For the problems of monotonicity and closeness we can do better.

2.1 Monotonicity

To detect whether a sequence is monotone or no, we can process the points while keeping a set of polar angles (field-of-view or FoV) that a direction vector v can have. Notice that the FoV is always an interval. When a new point arrives we can update the FoV in constant time. If the FoV ever becomes empty, then the given sequence is not monotone. This algorithm runs in $O(n)$ time for a sequence of length n . Using this submodule for *SN* we would get a $O(n^2 \log n)$ algorithm for computing $j^*(i)$ for all values of i .

However, doing a binary search for monotonicity is redundant, since we will do the same computation many times. Instead, for each i proceed with the

FoV computation until it is empty, which will happen precisely when we reach $j^*(i) + 1$. The resulting algorithm runs in $O(n^2)$ time.

2.2 Closeness

The simple method of closeness-check takes $O(n^2)$ for a sequence of n points, because one has to check all pairs of points. This results in SN runtime of $O(n^3 \log n)$, which is super-slow. We can improve the closeness-check to $O(n \log n)$ using furthest-point Voronoi diagrams, but that would not qualify for SN and would also take $O(n^2 \log^2 n)$ time to compute, which is also not fast.

Define $k^*(i)$ as the largest index such that $d(s_i, s_j) \leq 1$ for all $j \leq k^*(i)$. Clearly, $k^*(i) \geq j^*(i)$ because we are checking for a weaker property. Also, notice that k^* can be computed in $O(n^2)$ for all points by calculating all pairwise distances.

Noticing that

$$j^*(i) = \min(j^*(i+1), k^*(i)) \quad (1)$$

we can calculate $j^*(i)$ for all points in $O(n)$ extra time if we are given all values for $k^*(i)$. So we found a $O(n^2)$ algorithm for computing closeness property for all subsequences.

3 Monotonicity in Linear Time

Bokal et al. introduced an elegant framework to deal with range queries for hereditary problems. The key idea in their work is to greedily split the given set of points into ranges, defined by anchor points, and solve for each using a divide and conquer. To simplify the reasoning we first define a property matrix A_P such that

$$A_P(i, j) = \begin{cases} 1 & P(S[i, j]) = \text{true} \\ 0 & P(S[i, j]) = \text{false} \end{cases}$$

See Figure 3 for an example. Suppose that there is an algorithm J_P that, for a given i , finds $j^*(i)$ (notice that even if such J_P exists, we are not happy to run it n times for each i). We use J_P to find *anchor* points in S as follows. Let a_k be the index of k^{th} anchor point in S . We define $a_1 = s_1$, $a_k = \max(1 + a_{k-1}, j^*(a_{k-1}))$ for $k > 1$. Thus, the anchor points can be found using J_P . Figure 3 shows the anchor points on A (leftmost image).

Also shown in Figure 3 are rectangles formed using the anchor points. If we solve a rectangle (find the curve in A separating 1's from 0's) for each rectangle, then we will be done, as stated in Lemma 2.1 from Bokal et al.

Lemma 2.1 [Bokal et al.]: Assume that we have the following two sub-routines for sequence S and a property P :

- (a) For any $a = 1, \dots, n$ find the $j^*(a)$, taking $T_{greedy}(j^*(a) - a)$
- (b) Solve a rectangle with anchor points at opposite vertices (anchored rectangles). Suppose this module takes $T_{rect}(\text{height}(R) + \text{width}(R))$.




Figure 1: The property matrix A_P . The blue regions show the locations that are 1. The left image illustrates the anchor points. The middle one shows the greedily constructed rectangles from the anchor points. On the right side is the detailed view of the property matrix. Figure from [1].

Then we can solve the problem (i.e. find all $j^*(i)$) in time $O(n) + T_{greedy}(n) + T_{rect}(n)$.

Both of these modules are property specific, but Bokal et al. make another interesting generic step that explains how solving rectangles can be done using a divide-and-conquer approach. We will, however, skip that step and focus on the problem of monotonicity. It turns out that monotonicity does not require the divide-and-conquer approach to solve anchored rectangles.

3.1 Solving Monotonicity

Section 2.1 discusses how to find $j^*(i)$ in linear time for a given i by using intervals of polar angles. This algorithm will serve as the subroutine for Lemma 2.1(a).

Consider an anchored rectangle R with a lower-left corner (a, a) , width w and height h .

- By traversing the points $S[a, a+w]$, we can compute for each $j = a, \dots, a+w$ the interval of polar angles I_j such that $S[a, j]$ is monotone with respect to all directions in I_j .
- By traversing the points $S[a-h, a]$ in reverse order, a similar set I_i can be computed for $a-h \leq i < a$, representing the set of monotone directions for $S[i, a]$.

If $i < a < j$, then $S[i, j]$ is monotone if and only if I_i and I_j have a non-empty intersection. Thus, for each $i = a-h, \dots, a$, $j^*(i)$ is the largest index such that $I_i \cap I_{j^*(i)} \neq \emptyset$. Noting that search for $j^*(i+1)$ can start from $j^*(i)$, we can compute $j^*(i)$ for all $i \in [a-h, a]$ in linear time.

These two subroutines, together with Lemma 2.1 give a linear solution to the monotonicity problems. We have evaluated Naive algorithm (from Section 2) and this algorithm in Section 6.

4 Closeness in $O(n \log^2 n)$

Chan, Prat [2] take a different approach to calculating hereditary properties, they first efficiently calculate k^* values for all the pairs using a range tree [3] that has a secondary structure to make querying faster, and then they calculate j^* values from those.

During the first part of the algorithm we use a data structure that given a point i , could efficiently answer what is the point with smallest index $k^*(i) + 1$ that lies outside a unit circle centered at i . Notice that this is equivalent to checking whether s_i lines in the intersection of unit circles centered at points $s_{i+1}, \dots, s_{k^*(i)}$. The intersection of unit circles is convex and we call it *polyarcs*, as it is similar to a polygon, but has arcs of circles as edges instead of straight lines. In order to check whether a point is inside a polyarc formed by a range of points in the given sequence, we use a binary range tree.




Figure 2: A search path finding $p = k^*(q) + 1$ of point q . The search goes up until it find a range where p lies and then goes down narrowing down the range until it reaches the leaf p . Figure from [2].

First we build a regular 1 dimensional range tree using indices of the data points. Each node v of the range tree represents a continuous interval of a given points, and stores the polyarc that is formed by intersecting unit circles centered at these points.

Suppose we have an algorithm that can test whether a point is inside a polyarc in $Q(n)$ time. Then using the range tree we can find $k^*(i)$ for any point in $O(Q(n) \log n)$ time using the following algorithm. It starts the search from a leaf corresponding to the query point q (see Figure 4).

- In the first “up” phase, we walk upward from q towards the root. Each time we go up from a left child, we query the secondary structure at right sibling, to see if there exists a point in that range that is far from q .
 - If no, we can extend the search by going one more level up.
 - If yes, the first point far from q is in the right child, so we proceed to the second step to locate that node.

- In the second (“down”) phase, we walk downward from the current node to a leaf to locate the far point with smallest index. Each time we descend from a node, we query the secondary structure at the left child to see if there exists a point in the left child’s range that is far from q
 - If no, the point that’s far from q lies in the range of the right child, so we descend right.
 - If yes, we descend left.

After we find $k^*(i)$ for all the nodes, we combine those using Equation 1 in $O(n)$ time to find $j^*(i)$ for all i . This gives total runtime of $O(n \log n Q(n))$.

Now, to compute the intersection of polyarcs we use a relatively standard technique. It is possible to extend the sweep-line polygon intersection algorithm by Shamos and Hoey [4], for computing polyarc intersections in linear time, assuming that the vertices of the shapes are given in sorted order. We can construct the range-tree from bottom to up, at each level spending $O(n)$ time to intersect all polyarcs on that level. Therefore, the total construction time of a range tree is $O(n \log n)$.

By storing the vertices of the polyarcs in sorted order with respect to their x coordinate we can answer *contains* queries in $Q(n) = O(\log n)$ time using a binary search. This gives us an $O(n \log^2 n)$ algorithm to compute $j^*(i)$ for all i . The next section contains the details about the implementation of this algorithm.

5 Implementation Details

In the papers reviewed here ([1, 2]) there are at least 6 algorithms presented. We wanted to implement algorithm that use frameworks of both papers. In order to limit the size of the implementation we decided to pick monotonicity property for its simplicity (and elegance), which was solved in Bokal et al. [1]. We chose to implement the closeness testing algorithm from [2] because it used many ideas from the class, for example the intersection of convex envelopes in 2 dimensions, range trees.

The implementations of naive algorithms as well as of the efficient monotonicity testing algorithm are relatively easy, so in this paper we skip discussing them. We implemented the $O(n \log^2 n)$ algorithm for solving closeness testing from [2]. The algorithm has two distinct parts: a high-level tree-range data structure, and a lower-level data structure for storing *polyarcs* (see Section 4). Higher level algorithm was easier of the two, it involved implementing a range tree and the querying algorithm described in Section 4.

A polyarc is a convex geometric shape, which is defined by a clockwise set of its vertices like a polygon. Unlike a polygon, each edge of a polyarc is an arc of a circle. The algorithm needs to be able to find the intersection of two polyarcs, and also decide whether a query point is inside the polyarc.

Both of these algorithms are similar to those of polygons, but the majority of the implementation effort was spent to handle the cases specific to polyarcs. Our

representation of polyarcs has two “sides”: upper and lower, respectively the upper part of the complex shape and the lower one. To intersect two polyarcs we first find a minimal upper envelope i.e. for each x coordinate we recorded the polyarc that had a lower upper side. Similarly, we computed a maximal lower envelope. The area above the maximal lower envelope and minimal upper envelope is the intersection of two polyarcs. Each of these steps can be performed in linear time by a sweep line technique.

The envelope representation of polyarcs is also helpful for querying whether a point is inside the polyarc or no. One needs to simply do a binary search for the x coordinate of the candidate point. If a vertical line at that location crosses the higher envelope on a larger value of y than the lower envelope, then the query point is inside the polyarc.

Overall, we wrote more than 1200 lines of code in C++ for the implementation of the algorithms, and 800 lines of code for testing. Large part of the code is dedicated to the closeness algorithm in $O(n \log^2 n)$. We used C++ for the implementation so that the overhead that the language adds to the runtime is minimized. A simple language to write code in is Python, but we think that in this case the strong type checking of C++ helped us, because of a complex structure of the code.

6 Experimental Results

We have tested the performance of the algorithms we’ve implemented with synthetic datasets, that we have tried to model based on some potential problems where the algorithms could have been used.

6.1 Monotonicity

For monotonicity we have tried looking at a trajectory of a particle moving steadily along x axis, but with some noise added to the movement. And the goal is to find all subsequences of intervals that the trajectory had some constant direction. In Figure 3 you can see the property matrix (i.e. the region where the property holds) when the noise added is a Gaussian noise with standard deviations $\sigma_x = 0.05, \sigma_y = 1$.

You can see that the area of the region where the property is substantial, so we expect Naive algorithm to perform worse than the faster version that we’ve implemented. Figure ?? shows runtime differences in milliseconds between Naive algorithm and algorithm we presented from Bokal et al.

TODO: Informative plot for gaussian... Actually make plots double plots (side by side, for 2 examples we are going to show) TODO TODO

Whereas if we look at a different dataset, e.g. a random walk, the performance results could be different. Figure ?? shows the property matrix for a random walk, and Figure ?? shows that the runtimes of the two algorithms are actually very close in practice for this case.




Figure 3: Property matrix for movement monotonicity for time intervals for a point moving steadily in x direction with Gaussian noise with $\sigma_x = 0.05, \sigma_y = 1$.

points in dataset	Bokal et al.	Naive
1000	0.6	3.6
3333	2.1	12.8
10000	5.8	43.5
33333	19.3	142.2
100000	64.2	416.8
333333	202.0	1398.2

Figure 4: Runtimes (in milliseconds) of Naive and Bokal et. al for point moving in x direction with Gaussian noise. Mean value of $j^*(i) - i = 1520$

points in dataset	Bokal et al.	Naive
3000	1.9	2.6
10000	6.0	8.3
30000	18.0	25.0
100000	61.7	89.5
300000	190.5	258.3
1000000	629.5	909.6

Figure 5: Runtimes (in milliseconds) of Naive and Bokal et. al for random walk dataset. Mean value of $j^*(i) - i = 4.0$

6.2 Closeness

For testing the closeness algorithm, we have generated points from a random walk. At each step we uniform randomly pick a direction, and move 0.3 distance in that direction. And we want to know all subsequences of the walk that are clustered together within distance of 1.

Figure 6 shows the property matrix when running closeness algorithm on the random walk data, and ?? shows runtime comparisons for different dataset sizes.

We can see that 6 the area of the matrix where the property is true is significant, and that's why we are saving a lot of performance. If we were to look at the example from previous part, the difference would not be as big, Figures ?? and ?? demonstrate that. TODO: Question to Hayk: should we even show more results like this??

TODO(arsen): Figure out example when runtimes are close together.

TODO(arsen): Fix timing comparison

TODO: we picked datasets that do not have the worst case runtime, but are also not very trivial.




Figure 6: Property matrix for checking 1 distance closeness on a random walk with step size 0.3

7 Conclusion

TODO: Finish Conclusion after other sections are polished.

points in dataset	Chan Prat	Naive
1000	9.3	0.5
3333	28.3	1.5
10000	88.7	4.5
33333	307.1	15.5
100000	815.2	46.8
333333	2736.9	157.3

Figure 7: Runtimes (in milliseconds) of Naive and Chan, Prat on **moving point** dataset. Mean value of $j^*(i) - i = 0.16$

points in dataset	Chan/Prat	Naive
3000	79.3	1316.2
10000	269.7	5998.3
30000	886.3	26716.6
100000	2751.2	90934.6

Figure 8: Runtimes (in milliseconds) of Naive and Chan, Prat on **random walk** dataset. Mean value of $j^*(i) - i = 1013$

It was really challenging to do an implementation project in computational geometry, some sentences in the papers take 100+ lines of code to implement. Thus, for small datasets, added complexity is not worth the speedup gained by the algorithm.

This was a great project otherwise, we've implemented couple of algorithms that we have seen during the lecture - sweep line, range tree, intersection of convex polygons using envelopes. Also couple that we had not found in Bokal et al and Chan, Prat.

References

- [1] Drago Bokal, Sergio Cabello, and David Eppstein. Finding All Maximal Subsequences with Hereditary Properties. In Lars Arge and János Pach, editors, *31st International Symposium on Computational Geometry (SoCG 2015)*, volume 34 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 240–254, Dagstuhl, Germany, 2015. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.
- [2] Timothy M. Chan and Simon Pratt. Two Approaches to Building Time-Windowed Geometric Data Structures. In Sándor Fekete and Anna Lubiw, editors, *32nd International Symposium on Computational Geometry (SoCG 2016)*, volume 51 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 28:1–28:15, Dagstuhl, Germany, 2016. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik.

- [3] George S Lueker. A data structure for orthogonal range queries. In *Foundations of Computer Science, 1978., 19th Annual Symposium on*, pages 28–34. IEEE, 1978.
- [4] Michael Ian Shamos and Dan Hoey. Geometric intersection problems. In *Foundations of Computer Science, 1976., 17th Annual Symposium on*, pages 208–215. IEEE, 1976.